

TechEd 2013



Back to Indexes: The Original Culprit

Amit Bansal

(CTO, eDominator Systems, Peopleware India)
@A_Bansal





Image Source: <http://www.jeffcubos.com/2011/10/18/load-em-up/>

Takeaways...

Lessons you will learn

- Missing Index feature: When, Why, How?
- Left-based subset: is it so easy?
- Database Tuning Advisor: When, Why, How?
- Suspense 😊
- In a nutshell: The optimizer's choice of indexes !

Demo 1: Mistake No. 1

Missing Index feature

Missing Index feature

Lessons Learnt

- Please do not follow the recommendations 'blindly'
- Review before implementation
- What just happened in the demo was: a redundant index
 - (unless you have variety of selectivity)
- Keep an eye on unused indexes ! (redundant & duplicates)

Demo 2: Mistake No. 2

Left-based subset

Left-based subset

Lessons Learnt

- Understand the index key column order (very critical)
- Query filter happens using a Left-based subset mechanism
 - An index is 'fully seekable' by a query if all the 'predicate columns' of the query are also index key columns in a left-ordered fashion starting with the first column (one exception though)
- And that's why looking at the operator tool-tip is so important 😊

Demo 3: Mistake No. 3

Suspense 😊

Design an index...

```
SELECT C.ContactID, C.FirstName,  
C.EmailPromotion  
FROM Person.Contact2 AS C  
WHERE C.FirstName LIKE N'L%'           ← 4.3 % (872)  
      AND C.EmailPromotion = 1         ← 25.2 % (5044)  
      AND C.ContactID < 10000          ← 50.04 % (9994)  
OPTION (MAXDOP 1)
```

--returns 77 out of 19972 records

Design an index...

```
SELECT C.ContactID, C.FirstName, C.EmailPromotion
FROM Person.Contact2 AS C
WHERE C.FirstName LIKE 'L%'
      AND C.EmailPromotion = 1
      AND C.ContactID < 10000
OPTION (MAXDOP 1)
```

-- Option 1

```
CREATE INDEX ContactComposite4
ON Person.Contact2(FirstName, EmailPromotion)
```

-- Option 2

```
CREATE INDEX ContactComposite5
ON Person.Contact2(EmailPromotion, FirstName)
```

Design an index...

```
SELECT C.ContactID, C.FirstName, C.EmailPromotion
FROM Person.Contact2 AS C
WHERE C.FirstName LIKE 'L%'
      AND C.EmailPromotion = 1
      AND C.ContactID < 10000
OPTION (MAXDOP 1)
```

-- Option 1

```
CREATE INDEX ContactComposite4
ON Person.Contact2(FirstName, EmailPromotion)
```

-- Option 2

```
CREATE INDEX ContactComposite5
ON Person.Contact2(EmailPromotion, FirstName)
```



Let's ask
DTA

The Optimizer's choice of Indexes

Lessons Learnt

- Multi-column index:
 - The index can be used to seek on the second column if there is an equality predicate on the first column
 - True:
 - FirstName = 'L' AND EmailPromotion = 1
 - Partially True:
 - FirstName LIKE 'L%' AND EmailPromotion = 1
 - False
 - FirstName LIKE '%L' AND EmailPromotion = 1

The Optimizer's choice of Indexes

Lessons Learnt

- Single column index:
 - True:
 - FirstName LIKE 'L%'
 - EmailPromotion = 1
 - ContactID < 10000
 - False:
 - FirstName LIKE '%L'
 - ABS(EmailPromotion) = 1
 - ContactID + 1 < 10000

The Optimizer's choice of Indexes

Lessons Learnt

- Common guideline (as it is):
 - Most selective column should be the first column
- Common guideline (as it should be):
 - Most selective column should be the first column when all other column predicates use the equality operator
- SQL maintains HISTOGRAM only for the first column of the index

The Optimizer's choice of Indexes

Lessons Learnt

Predicate

[AdventureWorks].[Person].[contact2].[EmailPromotion] as [c].[EmailPromotion]=(1) AND [AdventureWorks].[Person].[contact2].[ContactID] as [c].[ContactID]<(10000) AND [AdventureWorks].[Person].[contact2].[FirstName] as [c].[FirstName] like N'L%'

Object

[AdventureWorks].[Person].[contact2].[ContactComposite4]
[c]

Output List

[AdventureWorks].[Person].[contact2].ContactID,
[AdventureWorks].[Person].[contact2].FirstName,
[AdventureWorks].[Person].[contact2].EmailPromotion

Seek Predicates

Seek Keys[1]: Start: [AdventureWorks].[Person].[contact2].FirstName, [AdventureWorks].[Person].[contact2].EmailPromotion >= Scalar Operator(N'L'), Scalar Operator((1)), End: [AdventureWorks].[Person].[contact2].FirstName < Scalar Operator(N'M')

Predicate

[AdventureWorks].[Person].[contact2].[ContactID] as [c].[ContactID]<(10000) AND [AdventureWorks].[Person].[contact2].[FirstName] as [c].[FirstName] like N'L%'

Object

[AdventureWorks].[Person].[contact2].[ContactComposite5]
[c]

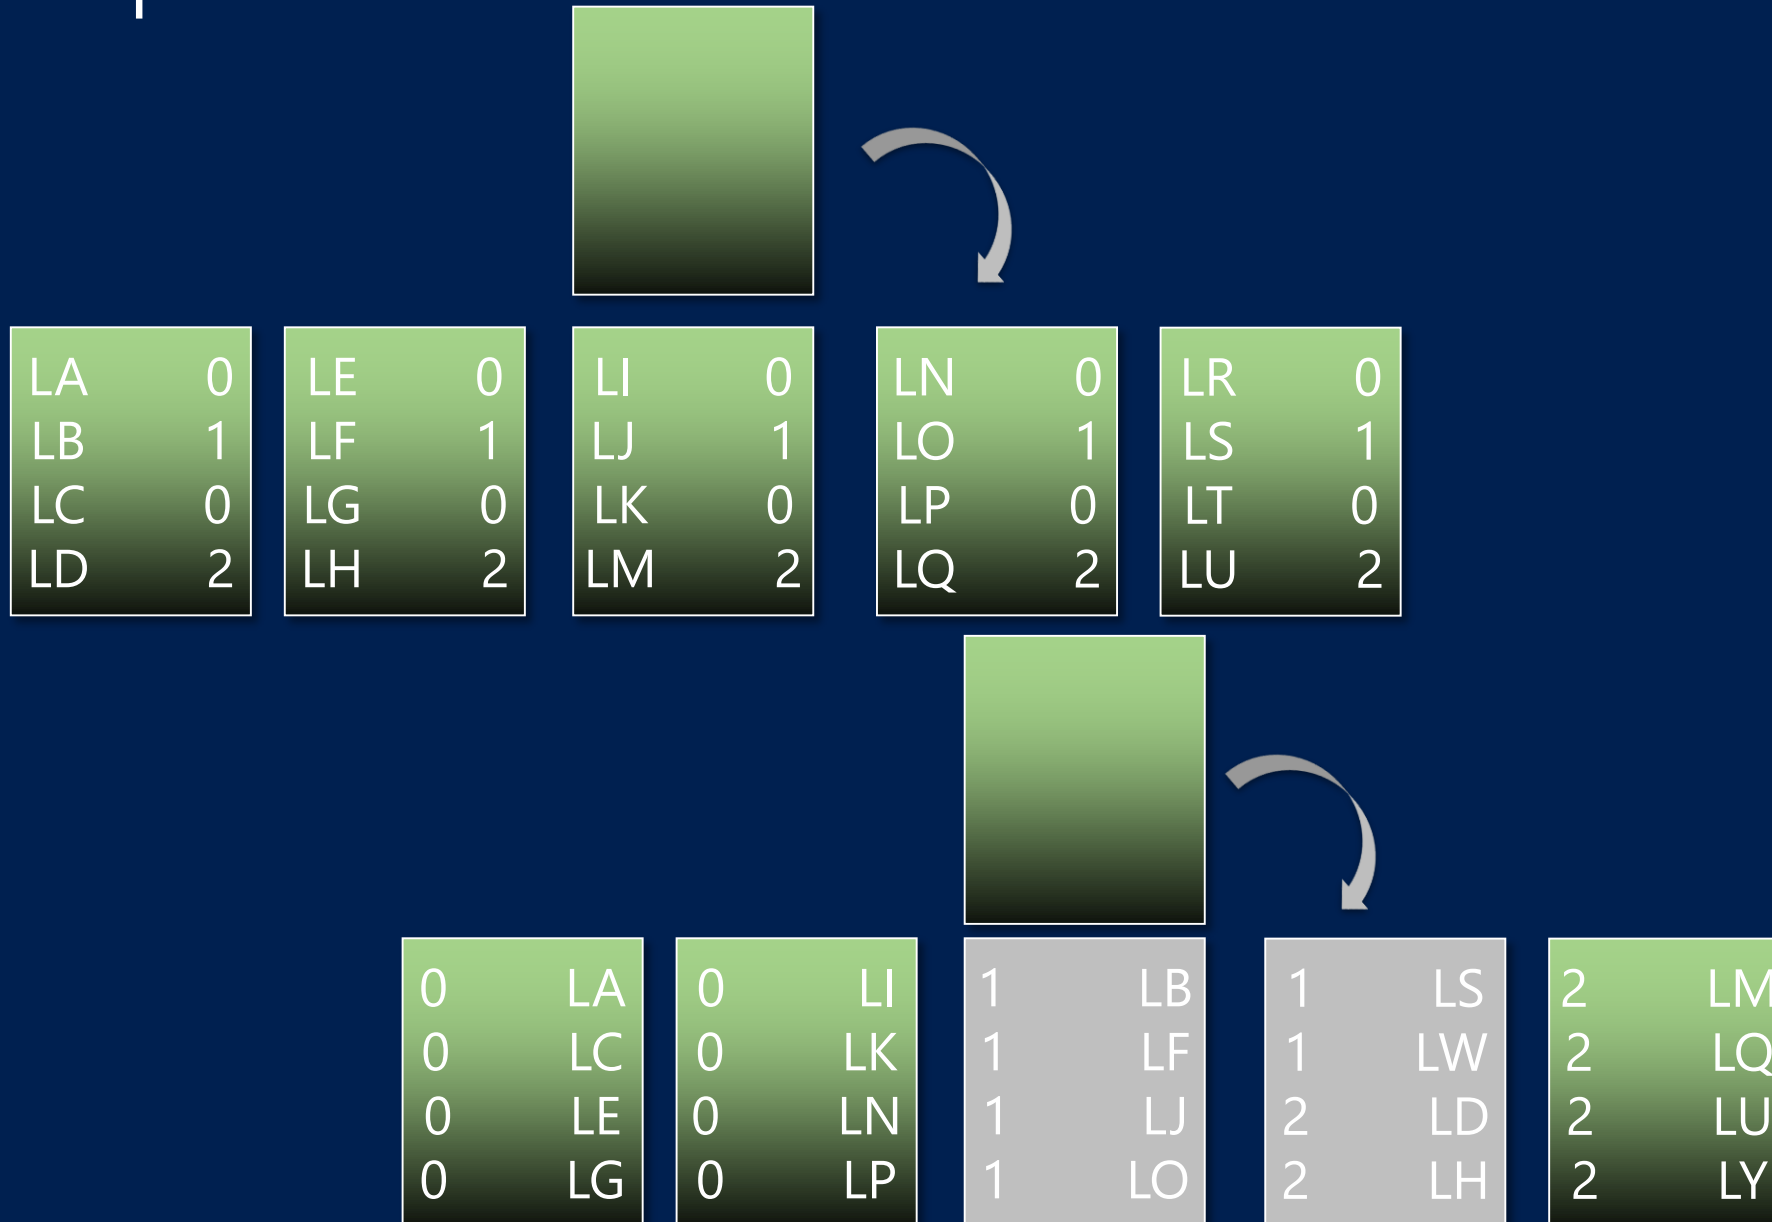
Output List

[AdventureWorks].[Person].[contact2].ContactID,
[AdventureWorks].[Person].[contact2].FirstName,
[AdventureWorks].[Person].[contact2].EmailPromotion

Seek Predicates

Seek Keys[1]: Prefix: [AdventureWorks].[Person].[contact2].EmailPromotion = Scalar Operator((1)), Start: [AdventureWorks].[Person].[contact2].FirstName >= Scalar Operator(N'L'), End: [AdventureWorks].[Person].[contact2].FirstName < Scalar Operator(N'M')

If time permits...



What about the DTA & Missing Index feature?



Sr. DBA



Jr. DBA



New to SQL Server

Takeaways...

Lessons you have learnt

- Missing Index feature: When, Why, How?
- Left-based subset: is it so easy?
- Database Tuning Advisor: When, Why, How?
- Suspense ☺ (Column Order in a Multi-column Index with equality & inequality operators)
- In a nutshell: The optimizer's choice of indexes !

Last, but not the least...

- There is no substitute to your deep knowledge about the optimizer and indexes. Period.

Summary/Call to Action

Follow me @A_Bansal

- Browse this recording once again 😊
- Download the slides & code snippets from www.SQLMaestros.com
- Try out the code snippets yourself
- Review your indexing strategies
- Implement the knowledge
- Try out various combinations in your 'test environment'

Evaluate this session

**Scan this QR
code to
evaluate this
session.**



Resources



Channel 9

Sessions on Demand

<http://channel9.msdn.com/Events/TechEd>

TechNet

Resources for IT Professionals

<http://microsoft.com/technet>

Learning

Microsoft Certification & Training Resources

www.microsoft.com/learning

msdn

Resources for Developers

<http://microsoft.com/msdn>

